
Thinking in Vigesimal, Answering in Decimal: How Language Models Represent French Numbers

Ari Holtzman and NeuriCo

Abstract

French uses a vigesimal (base-20) counting system for numbers 70–99, where expressions like *quatre-vingt-dix-sept* (literally “four-twenty-ten-seven”) encode implicit arithmetic operations ($4 \times 20 + 10 + 7 = 97$). We investigate how large language models represent and process these linguistically complex numerals through two complementary analyses: behavioral evaluation of GPT-4.1 on number conversion, arithmetic, and counting tasks, and representational analysis of MISTRAL-7B’s internal hidden states across standard French, Belgian French, and digit forms. At the behavioral level, GPT-4.1 achieves 100% accuracy on all tasks, including vigesimal numbers—refuting the hypothesis that irregular number words cause systematic errors in frontier models. However, probing MISTRAL-7B’s representations reveals that the vigesimal system leaves a distinctive internal fingerprint: 70% of vigesimal numbers’ nearest neighbors in representation space share a linguistic base (e.g., the *quatre-vingt* prefix) rather than numeric proximity. French number words are more linearly decodable ($R^2 = 0.979$) than digit strings ($R^2 = 0.943$), and Belgian French’s regular decimal system produces representations with the highest ordinal correlation ($\rho = 0.591$ vs. 0.520 for standard French). These findings demonstrate that models “think in vigesimal” even when they “answer in decimal,” revealing how linguistic structure shapes numerical representations beyond what behavioral benchmarks can detect.

1 Introduction

When a French speaker says *quatre-vingt-dix-sept*, they are not merely naming a number—they are performing implicit arithmetic: $4 \times 20 + 10 + 7 = 97$. This vigesimal (base-20) system, used for numbers 70–99 in standard French, stands in sharp contrast to the regular decimal systems of most European languages and even to Belgian French, which uses transparent forms like *nonante-sept* (ninety-seven). How do large language models (LLMs) handle this linguistic irregularity? Do they simply memorize the mapping, or does the vigesimal structure leave traces in their internal representations?

Understanding how LLMs bridge linguistic form and mathematical meaning is a central question in the study of numerical reasoning. Recent work has shown that models encode numbers using base-10 digit representations [Levy and Geva, 2024], that sinusoidal probes can recover numeric values from hidden states [Kadlcik et al., 2025, Stefanik et al., 2025], and that implicit compositional operators embedded in numeral systems pose a significant challenge [Bhattacharya et al., 2025]. However, no prior study has examined how the French vigesimal system—one of the most widely encountered irregular number systems—is represented internally by LLMs.

We address this gap with the first systematic study of French vigesimal number processing in LLMs. Our approach combines behavioral evaluation using GPT-4.1 with representational analysis of MISTRAL-7B’s hidden states, leveraging Belgian French as a natural control condition that isolates the effect of vigesimal structure from other linguistic factors.

Our behavioral results show that GPT-4.1 achieves perfect accuracy (100%) on all French number tasks, including vigesimal conversion, arithmetic, and counting across vigesimal boundaries. Yet beneath this behavioral ceiling, MISTRAL-7B’s internal representations tell a different story: vigesimal numbers cluster by linguistic base rather than numeric proximity, with 70% of nearest neighbors sharing a prefix like *quatre-vingt*. French number words yield higher linear probe accuracy ($R^2 = 0.979$) than digit strings ($R^2 = 0.943$), and Belgian French produces the most numerically ordered representations ($\rho = 0.591$).

We make the following contributions:

- We conduct the first behavioral and representational analysis of French vigesimal number processing in LLMs, demonstrating perfect behavioral accuracy alongside structured internal representations.
- We show that vigesimal numbers cluster by linguistic base in representation space, revealing that LLMs organize these numbers by their surface form rather than their numeric value.
- We establish Belgian French as an effective control condition, finding that its regular decimal system produces representations with superior ordinal structure compared to standard French.
- We find that French number words produce more linearly decodable numeric representations than raw digit strings, suggesting that linguistic context enriches rather than obscures numerical information.

2 Related Work

Number representations in LLMs. A growing body of work investigates how language models encode numerical information. Levy and Geva [2024] demonstrate that models use per-digit circular representations in base 10, achieving 91–96% probe accuracy for individual digits in Llama 3 and MISTRAL-7B. Kadlcik et al. [2025] introduce sinusoidal-basis probes that achieve near-perfect accuracy in decoding numeric values, overturning earlier findings of imprecise number representations. Stefanik et al. [2025] extend this work by showing that sinusoidal representations are universal across models, layers, and contexts, with cross-model Fourier frequency overlap of 1.0 for the top frequencies. Davies et al. [2025] caution that embedding-level representations are far from continuous scalar encodings, with string artifacts and tokenization effects dominating the embedding space. Marjeh et al. [2025] find that LLM number representations are entangled, blending string-level properties (edit distance) with numerical properties (log-linear magnitude). Our work extends this line of research to French vigesimal numbers, testing whether the base-10 internal representations documented in prior work accommodate numbers whose linguistic form encodes base-20 arithmetic.

Multilingual numerical reasoning. Bhattacharya et al. [2025] study cross-linguistic numeral puzzles and find that LLMs fail not because of mathematical difficulty but because they cannot infer implicit compositional operators (addition, multiplication) from numeral structure. Their work explicitly identifies French as an example vigesimal system, though they do not probe French representations directly. Rahman and Mishra [2025] demonstrate that LLM numerical reasoning is fragile—models succeed at deterministic tasks but fail at combinatorial puzzles, revealing pattern matching over genuine number sense. Unlike these studies, which focus on behavioral evaluation, we complement behavioral analysis with representational probing to understand *why* models succeed or fail.

Tokenization and arithmetic. Singh and Strouse [2024] show that tokenization direction significantly impacts arithmetic accuracy, with right-to-left tokenization improving performance by up to 20%. GraphPKU [2024] establish the NUPA benchmark with systematic evaluation of number representations and find that performance degrades with number length and that single-digit tokenizers outperform multi-digit ones. Various Authors [2026] analyze the mechanistic basis of counting in LLMs, finding logarithmic compression and specific attention heads that mediate count information. These findings are directly relevant to French vigesimal numbers, which require 40% more tokens than their decimal counterparts, creating additional computational overhead for the transformer.

Positioning of our work. Prior work on number representations has focused almost exclusively on English and digit tokens. We contribute the first study of French vigesimal representations, using Belgian French as a controlled comparison that isolates the effect of linguistic irregularity. Our combined behavioral-representational approach reveals structure invisible to behavioral benchmarks alone.

3 Methodology

We combine behavioral evaluation of a frontier model with representational analysis of an open-weight model. This two-pronged approach lets us assess both *what* models produce and *how* they internally organize numerical information.

3.1 Dataset Construction

No prior dataset exists for studying French number representations in LLMs. We construct a custom dataset covering numbers 0–999, with focused analysis on 0–99. Each entry contains the number, its standard French word form, a vigesimal flag, and the implicit arithmetic operations encoded in the word (if any). For the vigesimal range (70–99), we additionally include Belgian French forms (e.g., *septante* for 70, *huitante* for 80, *nonante* for 90), which use regular decimal constructions and serve as a controlled comparison.

We derive four sub-datasets from this data:

- **Number conversion** (100 items): French words and digit strings for numbers 0–99, used for bidirectional conversion tasks.
- **Belgian French subset** (30 items): Vigesimal numbers 70–99 with both standard and Belgian French forms.
- **Arithmetic tasks** (100 items): Addition problems with French operands, annotated for vigesimal involvement.
- **Counting prompts** (6 items): Boundary-crossing sequences at vigesimal transitions ($69 \rightarrow 70$, $79 \rightarrow 80$, $89 \rightarrow 90$) and matched decimal controls ($29 \rightarrow 30$, $39 \rightarrow 40$, $49 \rightarrow 50$).

3.2 Behavioral Evaluation

We evaluate GPT-4.1 on five task types with temperature set to 0 for deterministic outputs:

French-to-number conversion. Given a French number word (e.g., *quatre-vingt-treize*), the model must produce the corresponding digit (93). We test all 100 numbers from 0 to 99.

Number-to-French conversion. The inverse task: given a digit, produce the correct French word form.

Belgian French-to-number conversion. The same conversion task using Belgian French forms for the 30 vigesimal numbers.

Arithmetic in French. The model receives addition problems where operands are written as French words and must produce the sum. Of the 100 tasks, we annotate which involve at least one vigesimal operand or produce a vigesimal result.

Counting sequences. The model must continue a counting sequence in French across a boundary. We test three vigesimal boundaries ($69 \rightarrow 70$, $79 \rightarrow 80$, $89 \rightarrow 90$) and three decimal controls, evaluating whether each number in the generated sequence is correct.

3.3 Representational Analysis

We use MISTRAL-7B v0.3 to extract internal representations, as it is an open-weight model with strong French language capabilities. The model is loaded in float16 on a single NVIDIA RTX 3090.

Representation extraction. For each number $n \in \{0, \dots, 99\}$, we present three prompt forms: (1) *Le nombre est [French word]*, (2) *Le nombre est [digit string]*, and (3) *Le nombre est [Belgian French word]* (for $n \in \{70, \dots, 99\}$). We extract the hidden state at the last token position from five layers: the embedding layer (0), layers 8, 16, 24, and the final layer (32). This yields representation vectors $\mathbf{h}_n^{(\ell)} \in \mathbb{R}^d$ for each number n , form, and layer ℓ .

Tokenization analysis. We count the number of tokens produced by MISTRAL-7B’s tokenizer for each number form, comparing vigesimal French words (70–99) against decimal French words (20–69) and digit strings.

Table 1: Behavioral accuracy of GPT-4.1 on French number tasks. The model achieves perfect accuracy across all conditions, including vigesimal numbers (70–99) that require implicit base-20 arithmetic.

Task	Decimal (0–69)	Vigesimal (70–99)	Overall
French → Number	100%	100%	100%
Number → French	100%	100%	100%
Belgian French → Number	—	100%	100%
Arithmetic	100%	100%	100%
Counting sequences	100%	100%	100%

Table 2: Mean token counts for MISTRAL-7B’s tokenizer across number forms and ranges. Vigesimal French words require 40% more tokens than decimal French words.

Number Form	Decimal (20–69)	Vigesimal (70–99)
French word	5.2	7.3 (+40%)
Digit string	3.0	3.0

Linear probing. We train Ridge regression probes ($\alpha = 1.0$) to predict the numeric value n from the hidden-state vector $\mathbf{h}_n^{(\ell)}$. We report mean absolute error (MAE) and R^2 using 5-fold cross-validation, comparing performance across French, Belgian French, and digit forms, and between vigesimal and decimal subsets.

Representational Similarity Analysis. We compute pairwise cosine distances between all number representations within each form, producing a representational dissimilarity matrix (RDM). We compare each RDM against the ideal numeric RDM (where dissimilarity equals $|n_i - n_j|$) using Spearman correlation ρ , and compare RDMs across forms to measure cross-form alignment.

Nearest-neighbor analysis. For each number in the vigesimal range, we identify its nearest neighbor in representation space (by cosine similarity) and record (a) whether the neighbor shares the same linguistic base (e.g., both start with *quatre-vingt*) and (b) the numeric distance $|n - n_{\text{NN}}|$. We compare these statistics against the decimal range (20–69) using Mann-Whitney U tests.

4 Results

4.1 Behavioral Evaluation

Table 1 presents the behavioral results for GPT-4.1. The model achieves **100% accuracy** on every task, with no difference between vigesimal and decimal numbers. All six counting sequences—three crossing vigesimal boundaries and three decimal controls—produce perfectly correct French number words with correct transitions.

This result refutes the hypothesis that vigesimal numbers cause behavioral difficulties for frontier models. GPT-4.1 has fully internalized the French counting system from its training data, including all vigesimal irregularities.

4.2 Tokenization Overhead

Vigesimal French numbers require substantially more tokens than their decimal counterparts (table 2). The mean token count for vigesimal words is 7.3, compared to 5.2 for decimal French words—a 40% increase. At the extreme, *quatre-vingt-dix-sept* (97) requires 10 tokens, compared to just 3 for the digit string “97.” Digit representations are constant at 3.0 tokens regardless of vigesimal status. This increased token count means the model must propagate numeric information through more transformer steps for vigesimal numbers.

Table 3: Ordinal correlation (Spearman ρ) between representational distances and numeric distances at the final layer of MISTRAL-7B. Belgian French, with its regular decimal system, achieves the best ordinal structure. All $p < 10^{-280}$.

Form	Spearman ρ
French	0.520
Digit	0.485
Belgian French	0.591

Table 4: Linear probe results (Ridge regression, 5-fold CV) predicting numeric value from MISTRAL-7B final-layer hidden states. French words are more linearly decodable than digit strings. Best results in **bold**.

Form	MAE (overall)	R^2	Vigesimal MAE	Decimal MAE
French	2.72 ± 0.65	0.979	2.37	1.82
Belgian French	2.79 ± 0.53	0.977	2.54	1.77
Digit	4.18 ± 1.19	0.943	4.37	3.49

4.3 Ordinal Structure of Representations

We measure how well each form’s representational distances preserve numeric ordering using Spearman correlation between the RDM and the ideal numeric distance matrix (table 3). Belgian French representations achieve the highest ordinal correlation ($\rho = 0.591$, $p < 10^{-300}$), surpassing both standard French ($\rho = 0.520$) and, notably, digit representations ($\rho = 0.485$). The regular decimal system of Belgian French thus produces internal representations that are more numerically organized than either the vigesimal system or raw digits.

4.4 Linear Probe Accuracy

Table 4 shows Ridge regression probe results at the final layer. French number words produce the most linearly decodable representations ($R^2 = 0.979$, MAE = 2.72), followed closely by Belgian French ($R^2 = 0.977$, MAE = 2.79), with digit strings substantially behind ($R^2 = 0.943$, MAE = 4.18). Within each form, vigesimal numbers have higher MAE than decimal numbers, confirming that the vigesimal range is harder to decode linearly.

The finding that contextualized French words are more decodable than digits is counterintuitive. We attribute this to the prompt template “*Le nombre est [X]*,” which activates a numeric semantic frame, producing richer representations than raw digit strings.

4.5 Vigesimal Clustering by Linguistic Base

The central question of our representational analysis is whether vigesimal numbers cluster by linguistic base or numeric value. Table 5 presents the nearest-neighbor analysis for numbers 70–99 in French representation space.

70% of vigesimal numbers’ nearest neighbors share a linguistic base—e.g., both 83 (*quatre-vingt-trois*) and 93 (*quatre-vingt-treize*) begin with *quatre-vingt* and are representationally close despite being 10 apart numerically. Similarly, 73 (*soixante-treize*) and 63 (*soixante-trois*) cluster together because they share the *soixante*- prefix. The pattern is bimodal: many vigesimal numbers do follow numeric ordering (NN distance = 1), but a substantial fraction have nearest neighbors driven by shared prefixes.

A Mann-Whitney U test comparing vigesimal vs. decimal NN distances yields $U = 708.0$, $p = 0.667$ (not significant), indicating that vigesimal numbers are not *uniformly* worse—rather, the effect is concentrated in specific numbers where linguistic similarity overrides numeric proximity.

Table 5: Nearest-neighbor analysis for vigesimal French numbers (70–99). The majority of nearest neighbors share a linguistic base rather than being numerically adjacent. Examples show numbers whose representational nearest neighbor is numerically distant but linguistically similar.

Metric	Vigesimal (70–99)	Decimal (20–69)
NN shares linguistic base	70%	—
NN within ± 2 numerically	53.3%	34%
Mean NN error	4.3 ± 4.3	2.7 ± 1.7
Notable clustering examples:		
80 (<i>quatre-vingts</i>) $\rightarrow$ NN: 95 (<i>quatre-vingt-quinze</i>), $\Delta = 15$		
83 (<i>quatre-vingt-trois</i>) $\rightarrow$ NN: 93 (<i>quatre-vingt-treize</i>), $\Delta = 10$		
73 (<i>soixante-treize</i>) $\rightarrow$ NN: 63 (<i>soixante-trois</i>), $\Delta = 10$		
91 (<i>quatre-vingt-onze</i>) $\rightarrow$ NN: 82 (<i>quatre-vingt-deux</i>), $\Delta = 9$		

Table 6: Mean cosine similarity between representations of the same number across different forms in MISTRAL-7B’s final layer. French and Belgian French map to nearly identical internal representations despite different surface forms.

Comparison	Cosine Similarity
French–Digit	0.907
Belgian–Digit	0.907
French–Belgian	0.995

4.6 Cross-Form Similarity

Table 6 shows mean cosine similarity between representations of the same number across different forms. French and Belgian French representations are nearly identical (0.995 cosine similarity), suggesting the model maps both surface forms to a shared underlying representation. Both French forms are moderately aligned with digit representations (0.907).

4.7 Representational Similarity Analysis

Table 7 reports Spearman correlations between RDMs. Belgian French’s RDM is most correlated with the ideal numeric RDM ($\rho = 0.591$), confirming its superior ordinal structure. The French–Belgian RDM correlation ($\rho = 0.890$) is high but not as extreme as the per-number cosine similarity (0.995), indicating that while individual number representations are nearly identical, the *pairwise distance structure* differs—consistent with the linguistic clustering effect in the vigesimal range.

4.8 Layer Progression

Table 8 tracks how number representations develop across MISTRAL-7B’s layers. Probe accuracy improves throughout the network, with the best performance at the final layer (MAE = 15.0). The ordinal correlation peaks at early-to-middle layers ($\rho = 0.526$ at layer 8) before dipping in later layers and partially recovering. This pattern is consistent with the model processing linguistic features in early layers and reconstructing task-relevant numeric information in later layers.

5 Discussion

5.1 Behavioral Mastery vs. Representational Structure

The most striking finding is the disconnect between behavioral performance and internal representation. GPT-4.1 achieves 100% accuracy on all French number tasks, including vigesimal numbers—a result that refutes the initial hypothesis of 10–30% lower accuracy on vigesimal forms. This perfect performance likely reflects the sheer volume of French text in modern training corpora: the model has memorized the complete mapping between vigesimal words and their values.

Table 7: Representational Similarity Analysis (RSA): Spearman ρ between representational dissimilarity matrices (RDMs). Belgian French has the strongest alignment with ideal numeric ordering.

Comparison	Spearman ρ
French vs. Ideal Numeric	0.520
Digit vs. Ideal Numeric	0.485
Belgian vs. Ideal Numeric	0.591
French vs. Digit	0.656
Belgian vs. Digit	0.661
French vs. Belgian	0.890

Table 8: Layer-wise probe accuracy and ordinal correlation for French number representations in MISTRAL-7B. Numeric information develops progressively across layers.

Layer	Probe MAE	Spearman ρ
0 (embedding)	31.0	—
8	18.8	0.526
16 (middle)	16.4	0.525
24	18.4	0.490
32 (final)	15.0	0.520

Yet MISTRAL-7B’s internal representations reveal that this memorization preserves the linguistic structure of the vigesimal system. Numbers 80–99, which all share the *quatre-vingt* prefix, cluster together in representation space regardless of their numeric distance. The model “thinks in vigesimal”—organizing numbers by their linguistic base—even though it “answers in decimal” with perfect accuracy. This finding aligns with Marjeh et al. [2025]’s observation that LLM number representations entangle string-level and numerical properties, and extends it to a language where the string-number mapping is non-trivially compositional.

5.2 The Belgian French Advantage

Belgian French serves as a natural ablation of the vigesimal system. By replacing *soixante-dix* (60+10) with *septante* (70) and *quatre-vingt-dix* (4×20+10) with *nonante* (90), Belgian French removes the implicit arithmetic while preserving the language. We find that Belgian French produces representations with the highest ordinal correlation ($\rho = 0.591$ vs. 0.520 for standard French), confirming that **linguistic regularity directly improves the ordinal structure of numeric representations**.

The near-identical per-number cosine similarity between French and Belgian French (0.995) is equally notable. The model maps both *quatre-vingt-treize* and *nonante-trois* to essentially the same point in representation space, suggesting it recognizes them as the same number. However, the pairwise distance structure differs (RSA correlation = 0.890, not 1.0), because the vigesimal prefixes create distortions in the French RDM that are absent in the Belgian French RDM.

5.3 Why French Words Beat Digits in Probing

The counterintuitive finding that French words are more linearly decodable ($R^2 = 0.979$) than digit strings ($R^2 = 0.943$) deserves explanation. We hypothesize that the prompt template “*Le nombre est [X]*” activates a semantic frame that enriches the representation with numerical meaning. When the model processes *Le nombre est quarante-deux*, it has already been primed to interpret the following tokens as a number, producing a representation that is optimized for numeric content. In contrast, a bare digit string like “42” in the same template may be more ambiguous—it could be a code, an identifier, or a quantity—leading to a less numerically structured representation.

This result has practical implications: **linguistic context can improve rather than degrade numeric representations**, contrary to the assumption that words introduce noise relative to digits.

5.4 Limitations

Single model for representational analysis. We analyze only MISTRAL-7B for internal representations. The linguistic clustering effect may differ in models with different architectures, tokenizers, or training data distributions. Extending to Llama 3, Qwen, or other multilingual models would strengthen these findings.

Last-token extraction. We extract hidden states at the last token position only. Alternative aggregation strategies (mean pooling, attention-weighted pooling) might yield different representational structures, particularly for multi-token vigesimal words where information is distributed across tokens.

Context sensitivity. Our results are specific to the prompt template “*Le nombre est [X].*” Other contexts—arithmetic problems, counting sequences, free-form text—may produce different representational patterns.

Limited number range. We focus on 0–99. The vigesimal system also applies to hundreds (e.g., *deux-cent-quatre-vingt-dix-sept* = 297), where the compositional complexity increases further.

No causal intervention. We observe representational differences but do not establish causality. Activation patching [Levy and Geva, 2024] could test whether the linguistic clustering causally affects downstream outputs, but we leave this to future work.

6 Conclusion

We present the first systematic study of how LLMs represent and process French vigesimal numbers. Our behavioral evaluation shows that GPT-4.1 has fully mastered the French counting system, achieving 100% accuracy on conversion, arithmetic, and counting tasks across the vigesimal range. Our representational analysis of MISTRAL-7B reveals that this mastery coexists with a distinctive internal organization: vigesimal numbers cluster by linguistic base rather than numeric proximity, with 70% of nearest neighbors sharing a prefix. Belgian French’s regular decimal system produces representations with superior ordinal structure ($\rho = 0.591$ vs. 0.520), and French number words are more linearly decodable than digit strings ($R^2 = 0.979$ vs. 0.943).

These findings carry two broader implications. First, behavioral benchmarks alone are insufficient for understanding how models process numbers—probing reveals structure that perfect accuracy masks. Second, linguistic regularity directly shapes the quality of internal numeric representations, suggesting that the design of a language’s number system has consequences for how effectively LLMs encode numerical information.

Future work should extend this analysis to other models, longer number ranges, and causal interventions that test whether the vigesimal representational structure affects downstream reasoning. Investigating other languages with irregular number systems—such as Danish (which uses a different vigesimal system) or Georgian—would further clarify the relationship between linguistic form and numerical representation in LLMs.

References

- Antara Bhattacharya, Isabel Papadimitriou, Sam Davidson, and David Alvarez-Melis. Investigating the interaction of linguistic and mathematical reasoning using multilingual number puzzles. In *Proceedings of EMNLP*, 2025.
- Alex Davies et al. Language models do not embed numbers continuously. *arXiv preprint arXiv:2510.08009*, 2025.
- GraphPKU. Number cookbook: Number understanding of language models and how to improve it. *arXiv preprint arXiv:2411.03766*, 2024. ICLR 2025.
- Marek Kadlcik, Michal Stefanik, Timothee Mickus, Martin Spiegel, and Jaroslav Kuchar. Pre-trained language models learn remarkably accurate representations of numbers. In *Proceedings of EMNLP*, 2025.

- Amit Levy and Mor Geva. Language models encode numbers using digit representations in base 10. *arXiv preprint arXiv:2410.11781*, 2024.
- Raja Marjeh, Veniamin Veselovsky, Thomas L. Griffiths, and Ilia Sucholutsky. What is a number, that a large language model may know it? *arXiv preprint arXiv:2502.01540*, 2025.
- Mahir Rahman and Siddharth Mishra. A fragile number sense. *arXiv preprint arXiv:2509.06332*, 2025.
- Aarav Singh and DJ Strouse. Tokenization counts: The impact of tokenization on arithmetic in frontier LLMs. *arXiv preprint arXiv:2402.14903*, 2024.
- Michal Stefanik, Marek Kadlcik, Timothee Mickus, Martin Spiegel, and Jaroslav Kuchar. Unravelling the mechanisms of manipulating numbers in language models. *arXiv preprint arXiv:2510.26285*, 2025.
- Various Authors. Mechanistic interpretability of large-scale counting in LLMs. *arXiv preprint arXiv:2601.02989*, 2026.